



Petabyte Scale Data Warehousing Greenplum

Storage Considerations

Postgres Conf 2018

Marshall Presser

Craig Sylvester

Andreas Scherbaum

17 April 2018

Polymorphic Storage


Row-Oriented and Column-Oriented Tables

Row-oriented storage



- Supports mixed workloads (INSERT, UPDATE, DELETE, SELECT)
- Is supported with on both heap and append-optimized storage

Column-oriented storage



- Works well with data warehouse workloads
- Works well for data where you aggregate over a small number of columns
- Efficient for data where you modify a single column
- Supported on append-optimized storage

Table Storage Models



Customer	
PK	Customer_ID
	Customer_Name
	Customer_Street_Addr
	Customer_City
	Customer_State
	Customer_Country
	Customer_Postal_Code

Heap storage


- Default storage model
- Supports INSERT, UPDATE, DELETE
- Best for:
 - Data that is often modified
 - Smaller dimension tables
- Supports row-oriented tables
- Uses MVCC to support transactions

Append-optimized storage

- Append-optimized storage model:
- Optimized for data warehouses
- Works best with denormalized data
- Supports UPDATE and DELETE
- Best for:
 - Older data
 - Large fact tables
- Supports row and column-oriented tables
- Supports in-database compression
- Uses a Visibility Map (visimap) to hide outdated rows

Creating Heap and Append-Optimized Tables

Action	Example
Creating a heap, row-oriented table	<pre>CREATE TABLE tc_heap (id int, descr text) DISTRIBUTED BY (id);</pre>
Creating an append-optimized, row-oriented table	<pre>CREATE TABLE tc_ao (id int, sales float) WITH (appendonly=true) DISTRIBUTED BY (id);</pre>
Creating an append-optimized, column-oriented table	<pre>CREATE TABLE tc_ao_c (id int, sales float) WITH (appendonly=true, orientation=column) DISTRIBUTED BY (id);</pre>



Note: You cannot modify the storage or orientation of a table once defined. You can create a new table with the desired options and migrate your data.

Defining Append-Optimized Compression Tables

Action	Example
Creating a zlib compressed table with compression level 5	<pre>CREATE TABLE tc_ao_zlib5 (id int, sales float) WITH (appendonly=true, compresstype=zlib, compresslevel=5) DISTRIBUTED BY (id);</pre>
Creating a quicklz compressed table	<pre>CREATE TABLE tc_ao_quicklz (id int, sales float) WITH (appendonly=true, compresstype=quicklz) DISTRIBUTED BY (id);</pre>
Creating an AO table with an RLE compressed column and a zlib compressed column	<pre>CREATE TABLE tc_ao_rletype (id int, sales float ENCODING (compresstype=zlib, compresslevel=3), salesdate date ENCODING (compresstype=rle_type)) WITH (appendonly=true, orientation=column) DISTRIBUTED BY (id);</pre>

Compressing Table Data

Compression Algorithm	Compression Levels	Description	Table-Level Compression	Row-Level Compression
ZLIB	1 - 9	Offers the most compact ratio with a potential impact to CPU performance	Supported	Supported
QUICKLZ	1	Offers faster, but lower, data compression	Supported	Supported
RLE_TYPE Delta Compression (specific data types)	1 - 4	Offers run-length encoding compression for columns based on repeated values	Unsupported	Supported



Question: What type of data do you think would work well with the different offerings of compression?

Defining Default Table Storage Options

The diagram illustrates the hierarchy of `gp_default` storage options, ordered from highest to lowest priority. A large teal arrow points downwards, indicating the direction of decreasing priority.

Options	Level	Command
APPENDONLY	Object level	<code>CREATE TABLE ... WITH (...)</code>
BLOCKSIZE	Role level	<code>ALTER ROLE ... SET ...</code>
CHECKSUM	Database level	<code>ALTER DATABASE ... SET ...</code>
COMPRESSTYPE	System level	<code>gpconfig ...</code>
COMPRESSLEVEL		
ORIENTATION		

Highest priority (at the top) and **Lowest priority** (at the bottom) are indicated by the arrow.



Usage: Update default storage options at role level

```
names=> alter role student set
gp_default_storage_options='appendonly=true,compresstype=zlib';
Names=> set role student;
```

Additional Table Types



Temporary table

Temporary tables can be used for:

- Storing transient results needed for other session queries
- Perform transformations on data



Readable external table



Writable external table

External tables:

- Facilitate parallel data loading
- Stream data in from external sources
- Push data out of the database, in parallel

Temporary Tables – Overview

- Session-specific
- Dropped at the end of the session
- Take precedence over permanent tables of the same name
- Created in a special schema created on connection to a session
- Are distributed
- Can be indexed and analyzed

Temporary Tables – Two Use Cases

Customer	
PK	Customer_ID
	Customer_Name
	Customer_Street_Addr
	Customer_City
	Customer_State
	Customer_Country
	Customer_Postal_Code

**FILTE
R**

Customer	
PK	Customer_ID
	Customer_Name
	Customer_Street_Addr
	Customer_City
	Customer_State
	Customer_Country
	Customer_Postal_Code

Large table with billions of rows

Subset of larger table(s)

Working with smaller tables



Raw Source Data
(Serial Data Stream)



ETL Server

Raw Data - Block Loaded
(Parallel Data Streams - gpfdist)



**TRANSFORMATION
PROCESS**
(In Database Transforms
using SETs of data)

Data transformation

Creating a Temporary Table



Example: Creating a temporary

```
gpadmin=# CREATE TEMP[ORARY] TABLE monthlytranssummary (  
    storeid    INTEGER,  
    customerid INTEGER,  
    transmonth SMALLINT,  
    salesamttot DECIMAL(10,2)  
)  
ON COMMIT PRESERVE ROWS  
DISTRIBUTED BY (storeid, customerid);
```

You can

The following options to the `ON COMMIT` clause let you define how a temporary table is handled:

- `PRESERVE ROWS` – No action is taken on the table
- `DELETE ROWS` – The table is truncated
- `DROP` – The table is dropped