



Petabyte Scale Data Warehousing Greenplum

Postgres Conf 2018

Distribution

Marshall Presser
Craig Sylvester
Andreas Scherbaum
17 April 2018

CREATE TABLE Define Data Distributions

- One of the most important aspects of GP!
- Every table has a distribution method
- **DISTRIBUTED BY (column)**
 - Uses a hash distribution
- **DISTRIBUTED RANDOMLY**
 - Uses a random distribution which is not guaranteed to provide a perfectly even distribution
- Explicitly define a column or random distribution for all tables
 - Do not use the default

DISTRIBUTED BY (*column_name*)

- Use a single column that will distribute data across all segments evenly
- Example

```
CREATE TABLE foo (  
    id integer,  
    size float8)  
distributed by (id);
```
- For large tables significant performance gains can be obtained with local joins (co-located joins)
 - Distribute on the same column for tables commonly joined together
- Co-located join is performed within the segment
 - Segment operates independently of other segments
- Co-located join eliminates or minimizes motion operations
 - Broadcast motion or Redistribute motion

DISTRIBUTED RANDOMLY

- Uses a random algorithm
 - Distributes data across all segments
 - Minimal data skew but not guaranteed to have a perfectly even distribution
- Example

```
CREATE TABLE foo (  
    id integer,  
    size float8)  
distributed randomly;
```
- Any query that joins to a table that is distributed randomly will require a motion operation
 - Redistribute motion
 - Broadcast motion

DISTRIBUTED RANDOMLY

- Uses a random algorithm
 - Distributes data across all segments
 - Minimal data skew but not guaranteed to have a perfectly even distribution
- Example

```
CREATE TABLE foo (  
    id integer,  
    size float8)  
distributed randomly;
```
- Any query that joins to a table that is distributed randomly will require a motion operation
 - Redistribute motion
 - Broadcast motion

ALTERING DISTRIBUTION

- Distribution is done at table creation time, but can be altered.
 - `ALTER TABLE sales SET DISTRIBUTED BY (customer_id);`
- When you change the hash distribution of a table, table data is automatically redistributed.
- Changing the distribution policy to a random distribution does not cause the data to be redistributed.:
 - `ALTER TABLE sales SET DISTRIBUTED RANDOMLY;`
- Old school: create a new table with the new distribution and then swap.
 - `CREATE TABLE new_foo as select * from foo distributed randomly;`
 - `DROP TABLE foo;`
 - `ALTER TABLE new_foo RENAME to foo;`

DISTRIBUTED RANDOMLY

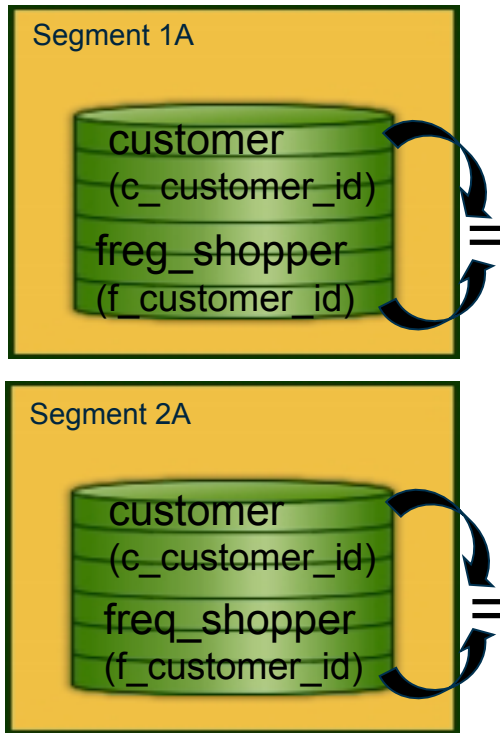
- Uses a random algorithm
 - Distributes data across all segments
 - Minimal data skew but not guaranteed to have a perfectly even distribution
- Example

```
CREATE TABLE foo (  
    id integer,  
    size float8)  
distributed randomly;
```
- Any query that joins to a table that is distributed randomly will require a motion operation
 - Redistribute motion
 - Broadcast motion

Hash Distributions: Data Skew and Computational Skew

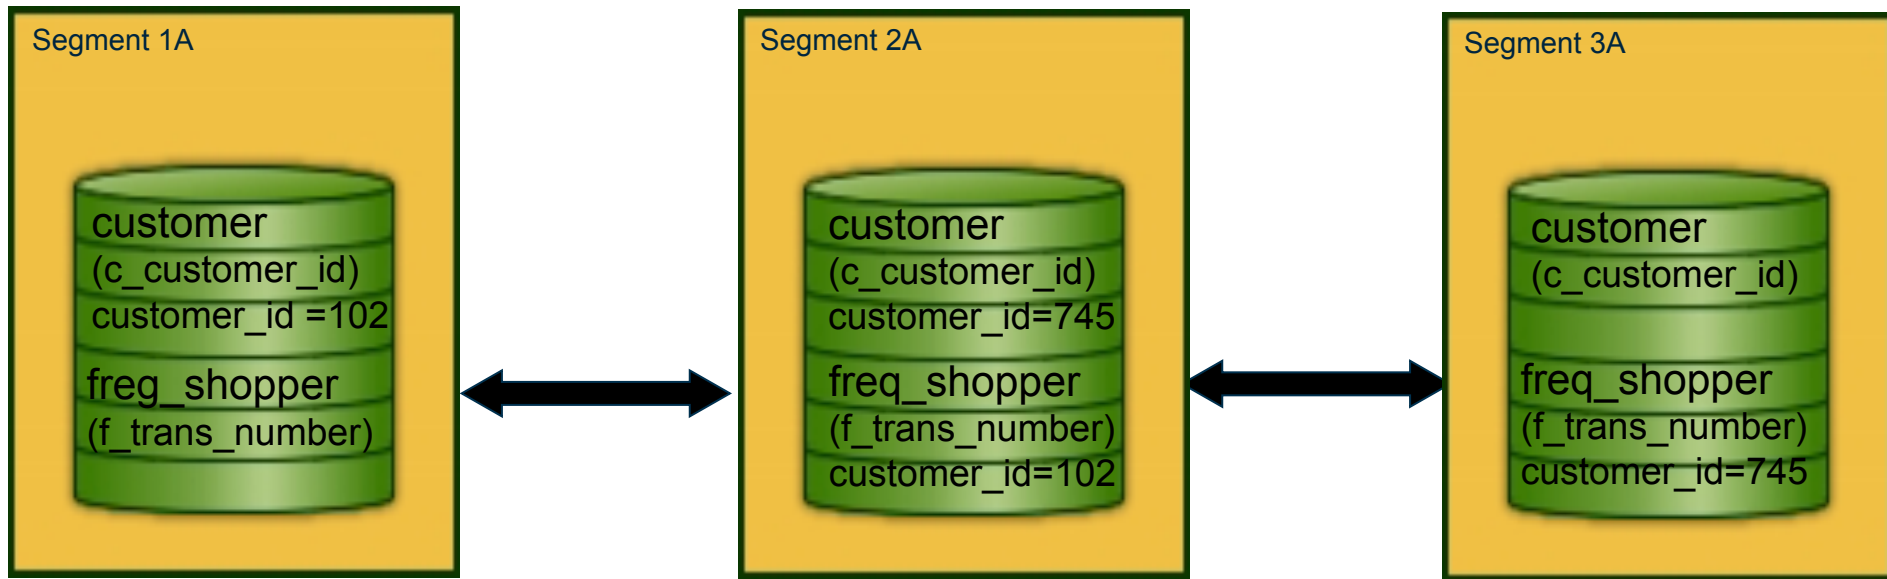
- Select a distribution key with unique values and high cardinality that will not result in data skew
 - Do not distribute on boolean keys and keys with low cardinality
 - The system distributes rows with the same hash value to the same segment instance therefore resulting in the data being located on only a few segments
- Select a distribution key that will not result in computational skew (in flight when a query is executing)
 - Operations on columns that have low cardinality or non-uniform distribution

Use the Same Distribution Key for Commonly Joined Tables




Distribute on the same key
used in the join
to obtain local joins

Redistribution Motion



WHERE customer.c_customer_id = freq_shopper.f_customer_id
freq_shopper table is dynamically redistributed on f_customer_id

Broadcast Motion



WHERE customer.c_statekey = state.s_statekey

The state table is dynamically broadcasted to all segments

Commonly Joined Large Tables Use the Same Data Type for Distribution Keys

<i>customer</i> (<i>c_customer_id</i>)	745::int
<i>freq_shopper</i> (<i>f_customer_id</i>)	745::varchar(10)

- Values might appear the same but they are stored differently at the disk level
- Values might appear the same but they HASH to different values
 - Resulting in *like* rows being stored on different segments
 - Requiring a redistribution before the tables can be joined

DISTRIBUTED BY (*column_name*)

- Do not distribute on columns that will be used in the WHERE clause of a query
- Do not distribute on dates or timestamps
- Never distribute and partition a table on the same column

Always Check for Data Skew on Initial Load and Subsequent Loads

- `SELECT COUNT(*), gp_segment_id FROM <table-name> GROUP BY gp_segment_id;`
- `SELECT 'facts' as "Table Name", max(c) as "Max Seg Rows", min(c) as "Min Seg Rows", (max(c)-min(c))*100.0/max(c) as "Percentage Difference Between Max & Min" from (SELECT count(*) c, gp_segment_id from facts group by 2) as a;`



gpuser=> SELECT count(*), gp_segment_id
from otp_x group by gp_segment_id order by 1;

count	gp_segment_id
-------	---------------

7800	7
------	---

699342	0
--------	---

1121978	5
---------	---

2002998	6
---------	---

2256304	2
---------	---

4125946	3
---------	---

5270874	1
---------	---

5374803	4
---------	---

--	--

(8 rows)

gpuser=> SELECT count(*), gp_segment_id from otp_r
group by gp_segment_id order by 1;

count	gp_segment_id
-------	---------------

2607504	7
---------	---

2607505	5
---------	---

2607505	4
---------	---

2607505	0
---------	---

2607505	2
---------	---

2607506	3
---------	---

2607507	1
---------	---

2607508	6
---------	---

--	--

(8 rows)